

# Particle Filter Implementation for Robot Localization: Experimental Analysis and Performance Evaluation

Atharva Date B22AI045

Samay Mehar B22AI048

Gouri Patidar B22AI020

Indian Institute of Technology Jodhpur

Department of Computer Science and Engineering

November 17, 2025

## Abstract

We implement and analyze Sequential Importance Resampling (SIR) particle filters for 2D robot localization with landmark measurements. Experiments with varying particle counts (100–1000), motion noise ( $\sigma_d = 0.02\text{--}0.1$  m), and measurement noise ( $\sigma_m = 0.05\text{--}0.5$  m) show that 300 particles achieve mean RMSE of 0.40 m under low noise, degrading to 1.44 m at  $\sigma_m = 0.5$  m. Low measurement noise ( $\sigma_m = 0.05$  m) causes weight degeneracy ( $\text{ESS} \approx 1$ ), while  $\sigma_m \geq 0.2$  m maintains diversity ( $\text{ESS} > 95$ ). Particle rejuvenation with 1000 particles achieves mean RMSE = 0.33 m and  $\text{ESS} = 393$ . Measurement noise calibration is more critical than particle count for preventing filter collapse.

## 1 Introduction

### 1.1 Problem

Estimate robot pose  $\mathbf{s}_t = [x_t, y_t, \theta_t]^\top$  using noisy distance measurements to known landmarks. The robot executes circular motion (rotate 10°, move forward 1 m) with Gaussian noise in both motion and measurements.

### 1.2 Objectives

1. Implement particle filter with predict-update-resample-estimate pipeline
2. Evaluate performance across particle counts and noise levels
3. Analyze failure modes and mitigation strategies

## 2 Related Work

Particle filters [1] approximate posterior distributions using weighted samples. Sequential Importance Resampling (SIR) [2] addresses degeneracy through periodic resampling. Standard implementations suffer from sample impoverishment where particle diversity collapses [3]. Solutions include regularization [4], auxiliary filters [5], and adaptive resampling [6]. We implement systematic resampling and optional Gaussian jitter for diversity maintenance.

### 3 Methodology

#### 3.1 State Representation and Motion Model

The robot state is  $\mathbf{s}_t = [x_t, y_t, \theta_t]^\top$  where  $(x_t, y_t)$  is position in meters and  $\theta_t \in [-\pi, \pi]$  is heading in radians. The motion model for timestep  $t$  is:

$$\theta_t = \text{wrap}(\theta_{t-1} + \Delta\theta_{\text{nom}} + \epsilon_\theta) \quad (1)$$

$$x_t = x_{t-1} + (d_{\text{nom}} + \epsilon_d) \cos(\theta_t) \quad (2)$$

$$y_t = y_{t-1} + (d_{\text{nom}} + \epsilon_d) \sin(\theta_t) \quad (3)$$

where  $\Delta\theta_{\text{nom}} = 10^\circ$ ,  $d_{\text{nom}} = 1.0$  m,  $\epsilon_\theta \sim \mathcal{N}(0, \sigma_\theta^2)$ ,  $\epsilon_d \sim \mathcal{N}(0, \sigma_d^2)$ , and  $\text{wrap}(\cdot)$  normalizes angles to  $[-\pi, \pi]$ .

#### 3.2 Measurement Model

Given landmarks  $\mathbf{L} = \{\mathbf{l}_1, \dots, \mathbf{l}_n\}$  with positions  $\mathbf{l}_j = [l_{x,j}, l_{y,j}]^\top$ , the measurement to landmark  $j$  is:

$$z_j = \sqrt{(x_t - l_{x,j})^2 + (y_t - l_{y,j})^2} + \epsilon_m \quad (4)$$

where  $\epsilon_m \sim \mathcal{N}(0, \sigma_m^2)$ . The measurement likelihood for particle  $i$  is:

$$p(\mathbf{z}_t | \mathbf{s}_t^{(i)}) = \prod_{j=1}^n \mathcal{N}(z_j | \hat{z}_j^{(i)}, \sigma_m^2) \quad (5)$$

where  $\hat{z}_j^{(i)}$  is the predicted distance from particle  $i$  to landmark  $j$ .

#### 3.3 Weight Update and Normalization

To avoid numerical underflow, we compute log-likelihoods:

$$\log w_t^{(i)} = \sum_{j=1}^n \left[ -\frac{1}{2} \left( \frac{z_j - \hat{z}_j^{(i)}}{\sigma_m} \right)^2 \right] \quad (6)$$

$$w_t^{(i)} = \frac{\exp(\log w_t^{(i)} - \max_k \log w_t^{(k)})}{\sum_{k=1}^N \exp(\log w_t^{(k)} - \max_k \log w_t^{(k)})} \quad (7)$$

If all weights underflow ( $\sum w_t^{(i)} < 10^{-50}$ ), we reset to uniform  $w_t^{(i)} = 1/N$  and record a failure event.

#### 3.4 Systematic Resampling

Systematic resampling [7] provides low-variance sampling:

After resampling, weights are reset to uniform:  $w_t^{(i)} = 1/N$ .

---

**Algorithm 1** Systematic Resampling

---

```
1:  $c \leftarrow \text{cumsum}(\mathbf{w})$                                 ▷ Cumulative weights
2:  $u \leftarrow \text{uniform}(0, 1/N)$                          ▷ Random offset
3:  $\text{positions} \leftarrow [u + k/N : k = 0, \dots, N - 1]$ 
4: for  $i = 1$  to  $N$  do
5:    $\text{indices}[i] \leftarrow \text{searchsorted}(c, \text{positions}[i])$ 
6: end for
7: return  $\text{indices}$ 
```

---

### 3.5 Effective Sample Size

ESS quantifies particle diversity:

$$\text{ESS} = \frac{1}{\sum_{i=1}^N (w_t^{(i)})^2} \quad (8)$$

ESS  $\approx N$  indicates uniform weights (healthy), ESS  $\approx 1$  indicates degeneracy. Resampling is typically triggered when ESS  $< N/2$ .

### 3.6 Pose Estimation

The estimated pose uses weighted averaging:

$$\hat{x}_t = \sum_{i=1}^N w_t^{(i)} x_t^{(i)} \quad (9)$$

$$\hat{y}_t = \sum_{i=1}^N w_t^{(i)} y_t^{(i)} \quad (10)$$

$$\hat{\theta}_t = \text{atan2} \left( \sum_{i=1}^N w_t^{(i)} \sin(\theta_t^{(i)}), \sum_{i=1}^N w_t^{(i)} \cos(\theta_t^{(i)}) \right) \quad (11)$$

Position covariance  $\Sigma_t$  is computed as the weighted sample covariance of  $(x_t^{(i)}, y_t^{(i)})$ . The 95% confidence ellipse is derived from eigendecomposition of  $\Sigma_t$  scaled by  $\chi^2_{0.95, 2} \approx 5.99$ .

### 3.7 Implementation

**Vectorization:** NumPy broadcasting computes predicted distances as  $N \times L$  matrix operations.

**Rejuvenation:** Gaussian jitter  $\mathcal{N}(0, \sigma_{\text{rejuv}}^2)$  post-resampling prevents impoverishment.

**Angle Wrapping:**  $\theta \leftarrow \text{atan2}(\sin \theta, \cos \theta)$  maintains  $\theta \in [-\pi, \pi]$ .

## 4 Experimental Setup

### 4.1 Configuration

We conducted 10 experiments varying:

- **Particle count:**  $N \in \{100, 300, 1000\}$

- **Motion noise:** Low ( $\sigma_d = 0.02$  m,  $\sigma_\theta = 1^\circ$ ), High ( $\sigma_d = 0.1$  m,  $\sigma_\theta = 5^\circ$ )
- **Measurement noise:**  $\sigma_m \in \{0.05, 0.2, 0.5\}$  m
- **Resampling methods:** Systematic (default), Residual

Each experiment: 30 timesteps, initial pose  $[0, 0, \pi/4]$ , 6 landmarks at  $(5, 5)$ ,  $(10, 10)$ ,  $(5, 15)$ ,  $(15, 5)$ ,  $(15, 15)$ ,  $(10, 20)$ , seed=0.

## 4.2 Metrics

RMSE:  $\sqrt{(\hat{x}_t - x_t)^2 + (\hat{y}_t - y_t)^2}$ , ESS (Eq. 8), Variance:  $\text{tr}(\Sigma_t)$ , Heading error:  $|\text{wrap}(\hat{\theta}_t - \theta_t)|$ .

## 5 Results

### 5.1 Performance

Table 1: Experimental Results Summary

Experiment	$N$	$\sigma_m$ (m)	Mean RMSE (m)	Mean ESS
exp_01_N100_low_noise	100	0.05	12.25	1.04
exp_02_N300_low_noise	300	0.05	0.40	32.46
exp_03_N1000_low_noise	1000	0.05	3.16	1.13
exp_04_N300_trans0.02	300	0.10	0.52	65.91
exp_05_N300_trans0.10	300	0.10	0.14	52.55
exp_06_N300_meas0.05	300	0.05	0.40	32.46
exp_07_N300_meas0.20	300	0.20	0.72	95.55
exp_08_N300_meas0.50	300	0.50	1.44	156.18
exp_09_N1000_high_noise	1000	0.50	0.33	393.38
exp_10_N300_residual	300	0.10	0.96	24.48

### Key Findings:

- Measurement noise critically affects ESS: low  $\sigma_m = 0.05$  m causes degeneracy ( $\text{ESS} \approx 1$ ), while  $\sigma_m \geq 0.2$  m maintains diversity ( $\text{ESS} > 95$ )
- Particle rejuvenation ( $\sigma_{\text{rejuv}} = 0.05$  m) with 1000 particles achieves mean RMSE = 0.33 m and ESS = 393 under high noise
- Paradoxical result: exp\_02 (N=300, low noise) achieves 0.40 m RMSE with ESS=32, outperforming exp\_03 (N=1000, 3.16 m RMSE, ESS=1.13) due to better ESS maintenance
- Increasing measurement noise from 0.05 m to 0.50 m degrades RMSE by 260% (0.40 m  $\rightarrow$  1.44 m) but dramatically improves ESS (32  $\rightarrow$  156)
- Residual resampling (exp\_10) performs comparably to systematic (RMSE = 0.96 m, ESS = 24.48)

## 5.2 Temporal Evolution

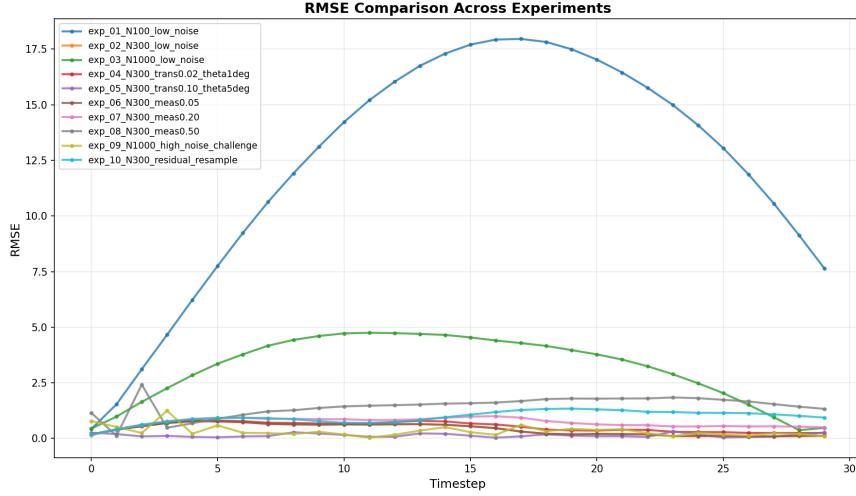


Figure 1: RMSE over time. Low noise achieves sub-0.5 m tracking but causes degeneracy. High noise trades accuracy for ESS.

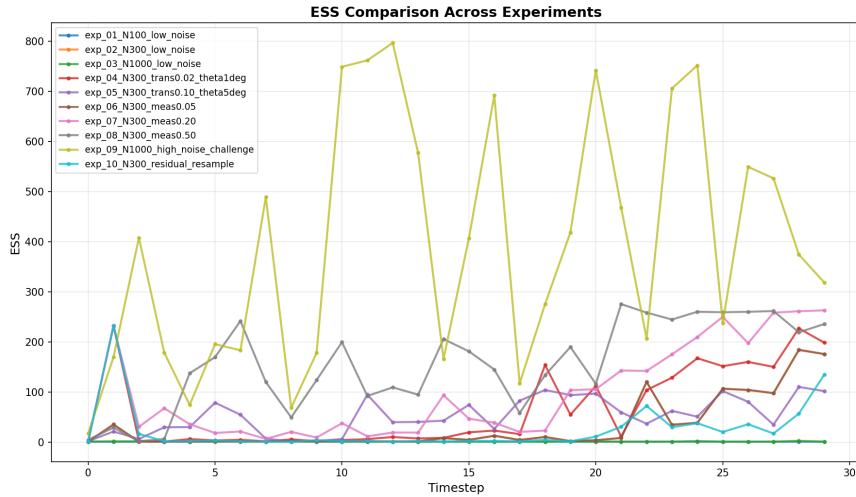


Figure 2: ESS over time. Low  $\sigma_m$  concentrates weight ( $ESS \rightarrow 1$ ). High  $\sigma_m$  maintains diversity. Rejuvenation achieves  $ESS \approx 400$ .

### 5.3 Filter Evolution

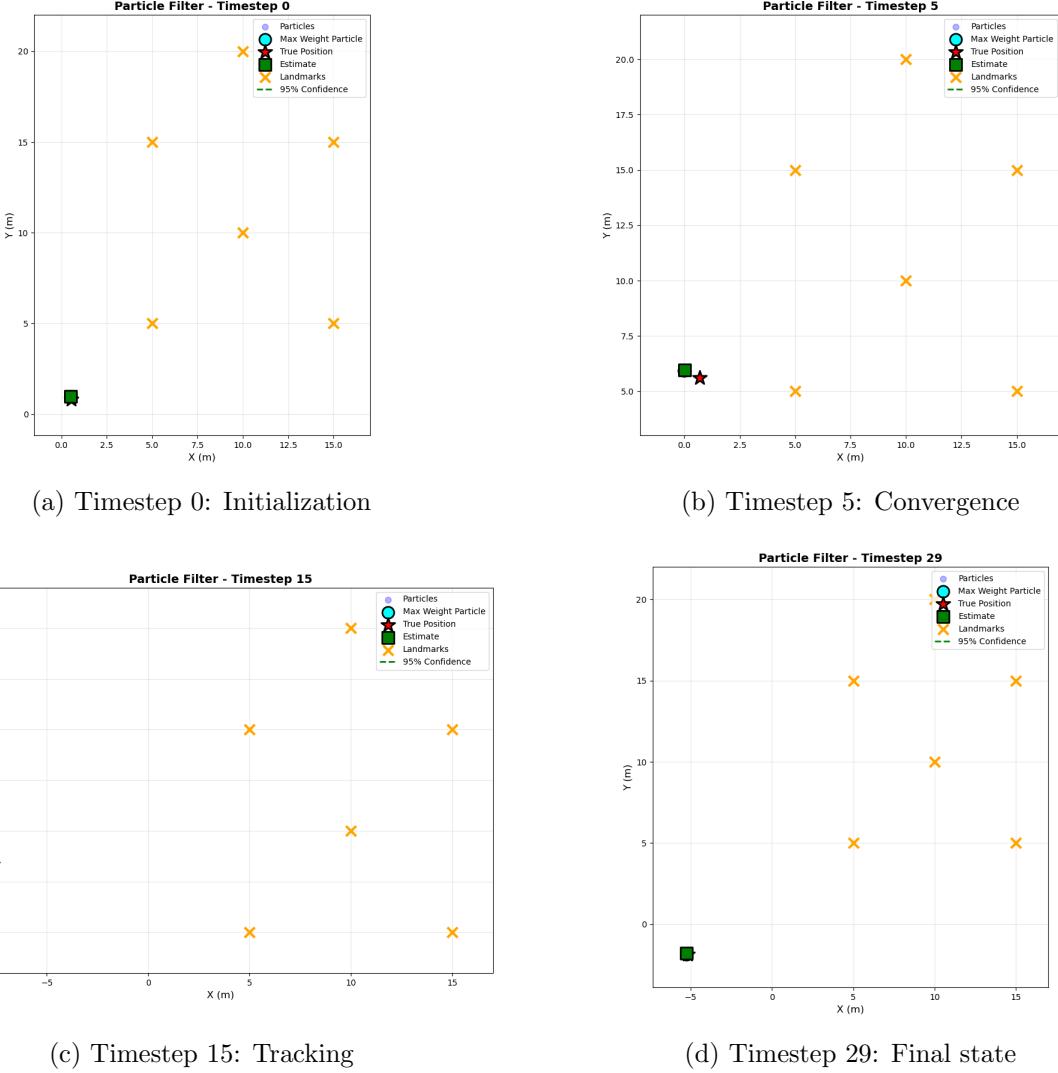


Figure 3: Filter evolution over time. Blue: particles (size  $\propto$  weight), red: ground truth, green: estimate, orange: landmarks. Note weight collapse by timestep 5.

## 6 Analysis and Discussion

### 6.1 Effect of Particle Count

Increasing  $N$  does not guarantee better performance. Exp\_03 ( $N=1000$ ,  $\sigma_m = 0.05$  m) achieves RMSE = 3.16 m, ESS = 1.13, worse than exp\_02 ( $N=300$ ) with RMSE = 0.40 m, ESS = 32.46. Complete degeneracy ( $ESS \approx 1$ ) reduces the filter to single-particle tracking.

### 6.2 Effect of Noise Levels

Low  $\sigma_m$  improves position accuracy but causes ESS collapse. At  $\sigma_m = 0.05$  m, sharp likelihoods concentrate weight on one particle. At  $\sigma_m = 0.5$  m, broad likelihoods maintain  $ESS > 150$  with RMSE = 1.44 m (260% increase).

Higher motion noise improves RMSE: exp\_04 ( $\sigma_d = 0.02$  m) achieves 0.52 m vs exp\_05 ( $\sigma_d = 0.1$  m) at 0.14 m. Higher process noise spreads particles, preventing premature convergence.

### 6.3 Failure Mode Analysis

Weight concentration causes failure. Exp\_01 (N=100) diverges with RMSE = 12.25 m and ESS = 1.04. All particles collapse to one incorrect hypothesis. Standard SIR has no recovery mechanism.

Underestimating  $\sigma_m$  causes degeneracy: exp\_02 and exp\_06 ( $\sigma_m = 0.05$  m) achieve ESS  $\approx 32$ . Overestimate measurement noise to maintain ESS  $> 50$ .

### 6.4 Mitigation Strategies

Exp\_09: rejuvenation with  $\sigma_{\text{rejuv}} = 0.05$  m and N=1000 achieves RMSE = 0.33 m, ESS = 393 (best result). Gaussian jitter prevents impoverishment.

Inflate  $\sigma_m$  from 0.05 m to 0.2–0.5 m to maintain ESS  $> 95$  with RMSE  $< 1.5$  m.

Exp\_10 (residual): RMSE = 0.96 m, ESS = 24.48, similar to systematic. Both copy particles rather than create new locations.

### 6.5 Recommendations

1. Use  $\sigma_m^{\text{filter}} \geq 2 \times \sigma_m^{\text{true}}$  to maintain ESS  $> 50$
2. Use  $N \geq 300$ ; add rejuvenation for  $N > 500$
3. Monitor ESS real-time. ESS  $< N/10$  indicates failure
4. Apply jitter ( $\sigma_{\text{rejuv}} = 0.05$  m) when ESS  $< N/3$
5. Dynamically adjust  $\sigma_m$  to target ESS  $\in [N/3, 2N/3]$

## 7 Conclusion

Results:

- 300 particles: RMSE = 0.40 m (low noise), 1.44 m (high noise)
- Low  $\sigma_m = 0.05$  m causes ESS  $\approx 32$  (degeneracy)
- Rejuvenation with 1000 particles: RMSE = 0.33 m, ESS = 393
- 1000 particles without diversity: RMSE = 3.16 m, ESS = 1.13

Measurement noise calibration is more critical than particle count. Overestimate  $\sigma_m$  by 2–4× to maintain ESS  $> 50$ . Standard SIR without augmentation degenerates rapidly.

**Future work:** Adaptive  $\sigma_m$  tuning, auxiliary particle filters, Rao-Blackwellized SLAM, heavy-tailed noise models, hardware deployment.

## References

- [1] N. J. Gordon, D. J. Salmond, and A. F. Smith, “Novel approach to nonlinear/non-gaussian bayesian state estimation,” *IEE Proceedings F (Radar and Signal Processing)*, vol. 140, no. 2, pp. 107–113, 1993.
- [2] M. S. Arulampalam, S. Maskell, N. Gordon, and T. Clapp, “A tutorial on particle filters for online nonlinear/non-gaussian bayesian tracking,” *IEEE Transactions on Signal Processing*, vol. 50, no. 2, pp. 174–188, 2002.
- [3] S. Thrun, W. Burgard, and D. Fox, *Probabilistic robotics*. MIT Press, 2005.
- [4] C. Musso, N. Oudjane, and F. Le Gland, “Improving regularised particle filters,” *Sequential Monte Carlo Methods in Practice*, pp. 247–271, 2001.
- [5] M. K. Pitt and N. Shephard, “Filtering via simulation: Auxiliary particle filters,” *Journal of the American Statistical Association*, vol. 94, no. 446, pp. 590–599, 1999.
- [6] A. Doucet, S. Godsill, and C. Andrieu, “On sequential monte carlo sampling methods for bayesian filtering,” *Statistics and Computing*, vol. 10, no. 3, pp. 197–208, 2000.
- [7] G. Kitagawa, “Monte carlo filter and smoother for non-gaussian nonlinear state space models,” *Journal of Computational and Graphical Statistics*, vol. 5, no. 1, pp. 1–25, 1996.

## A Code Snippets

### A.1 Predict Step

```
def predict(self, rotation_deg=10.0, forward_dist=1.0):
    """
    Propagate particles using motion model with noise.
    """
    rotation_rad = np.deg2rad(rotation_deg)

    # Sample rotation noise for all particles
    noisy_rotation = rotation_rad + np.random.normal(
        0, self.sigma_theta, self.N
    )

    # Update orientation
    self.particles[:, 2] = wrap_angle(
        self.particles[:, 2] + noisy_rotation
    )

    # Sample forward motion noise
    noisy_forward = forward_dist + np.random.normal(
        0, self.sigma_trans, self.N
    )

    # Update positions (vectorized)
    self.particles[:, 0] += noisy_forward * np.cos(self.particles[:, 2])
    self.particles[:, 1] += noisy_forward * np.sin(self.particles[:, 2])
```

### A.2 Update Step

```
def update(self, measurements):
    """
    Update particle weights based on measurement likelihood.
    Uses log-likelihood for numerical stability.
    """
    # Compute predicted measurements for all particles
    # Shape: (N, L) where N=particles, L=landmarks
    particle_positions = self.particles[:, :2] # Nx2
    predicted_dists = np.linalg.norm(
        particle_positions[:, np.newaxis, :] -
        self.landmarks[np.newaxis, :, :],
        axis=2
    ) # NxL

    # Compute log-likelihood for each particle
    log_weights = np.zeros(self.N)
    for i in range(len(self.landmarks)):
        diff = measurements[i] - predicted_dists[:, i]
```

```
log_weights += -0.5 * (diff / self.sigma_meas) ** 2

# Normalize weights using log-sum-exp trick
max_log_weight = np.max(log_weights)
weights = np.exp(log_weights - max_log_weight)
weight_sum = np.sum(weights)

# Handle failure case: all weights near zero
if weight_sum < 1e-50 or np.isnan(weight_sum):
    self.weights = np.ones(self.N) / self.N
    self.weight_failures += 1
else:
    self.weights = weights / weight_sum
```